

March 1, 2020

## 1 Assignment 5

1. Choose a regression dataset (bikeshare is allowed), perform a test/train split, and build a regression model (just like in assignment 3), and calculate the
  - Training Error (MSE, MAE)
  - Testing Error (MSE, MAE)
2. Choose a classification dataset (not the adult.data set, The UCI repository has many datasets as well as Kaggle), perform test/train split and create a classification model (your choice but DecisionTree is fine). Calculate
  - Accuracy
  - Confusion Matrix
  - Classification Report
3. (Bonus) See if you can improve the classification model's performance with any tricks you can think of (modify features, remove features, polynomial features)

```
[114]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import PolynomialFeatures
from sklearn.model_selection import train_test_split
from sklearn.linear_model import Ridge
from sklearn.metrics import (accuracy_score,
                             classification_report,
                             confusion_matrix, auc, roc_curve,
                             mean_squared_error)
from sklearn.tree import DecisionTreeClassifier
```

## 2 Regression Data Set

```
[30]: credit = pd.read_csv('../data/Credit.csv')
credit
```

```
[30]: Unnamed: 0    Income  Limit  Rating  Cards  Age  Education  Gender  \
0              1    14.891   3606    283     2    34           11    Male
1              2   106.025   6645    483     3    82           15   Female
```

2	3	104.593	7075	514	4	71	11	Male
3	4	148.924	9504	681	3	36	11	Female
4	5	55.882	4897	357	2	68	16	Male
..	...	...	...	...	...	...		
395	396	12.096	4100	307	3	32	13	Male
396	397	13.364	3838	296	5	65	17	Male
397	398	57.872	4171	321	5	67	12	Female
398	399	37.728	2525	192	1	44	13	Male
399	400	18.701	5524	415	5	64	7	Female

	Student	Married	Ethnicity	Balance
0	No	Yes	Caucasian	333
1	Yes	Yes	Asian	903
2	No	No	Asian	580
3	No	No	Asian	964
4	No	Yes	Caucasian	331
..	...	...	...	...
395	No	Yes	Caucasian	560
396	No	No	African American	480
397	No	Yes	Caucasian	138
398	No	Yes	Caucasian	0
399	No	No	Asian	966

[400 rows x 12 columns]

### 3 Classification Data Set

```
[78]: cancer_data = pd.read_csv('../data/breast-cancer-wisconsin.data')
```

```
[79]: cancer_data.columns = ['id',
                             'clump_thickness',
                             'cell_size',
                             'cell_shape',
                             'adhesion',
                             'epithelial_cell_size',
                             'bare_nucleoli',
                             'bland_chromatin',
                             'normal_nucleoli',
                             'mitosis',
                             'cell_class']
```

```
[80]: cancer_data
```

```
[80]:
```

	id	clump_thickness	cell_size	cell_shape	adhesion	\
0	1002945	5	4	4	5	
1	1015425	3	1	1	1	

2	1016277	6	8	8	1
3	1017023	4	1	1	3
4	1017122	8	10	10	8
..	...	...	...	...	...
693	776715	3	1	1	1
694	841769	2	1	1	1
695	888820	5	10	10	3
696	897471	4	8	6	4
697	897471	4	8	8	5

  

	epithelial_cell_size	bare_nucleoli	bland_chromatin	normal_nucleoli	\
0	7	10	3	2	
1	2	2	3	1	
2	3	4	3	7	
3	2	1	3	1	
4	7	10	9	7	
..	...	...	...	...	
693	3	2	1	1	
694	2	1	1	1	
695	7	3	8	10	
696	3	4	10	6	
697	4	5	10	4	

  

	mitosis	cell_class
0	1	2
1	1	2
2	1	2
3	1	2
4	1	4
..	...	...
693	1	2
694	1	2
695	2	4
696	1	4
697	1	4

[698 rows x 11 columns]

## 4 Question 1

Choose a regression dataset (bikeshare is allowed), perform a test/train split, and build a regression model (just like in assignment 3), and calculate the + Training Error (MSE, MAE) + Testing Error (MSE, MAE)

```
[101]: # separate groups from the data and only keep numerical data.
numerical_cats = credit.select_dtypes(include=['int64', 'float64'])
```

```
x_train, x_test, y_train, y_test = train_test_split(numerical_cats.
↳drop(['Rating'], axis=1),numerical_cats.Rating, test_size=.20)
x_train,y_train
```

```
[101]: (      Unnamed: 0   Income  Limit  Cards  Age  Education  Balance
      368      369   89.000   5759     3   37         6      345
      23       24   64.027   5179     5   48         8      411
     336      337   32.856   5884     4   68        13      926
     201      202   73.914   7333     6   67        15     1048
      36       37   62.413   6457     2   71        11      762
      ..      ...      ...      ...      ...      ...      ...
     322      323   27.229   3484     6   51        11      265
     378      379   19.349   4941     1   33        19      717
     339      340  149.316  10278     1   80        16     1107
      53       54   16.304   5466     4   66        10      957
     112      113   46.007   6637     4   42        14     1046
```

```
[320 rows x 7 columns], 368      440
      23      398
     336      438
     201      529
      36      455
      ...
     322      282
     378      366
     339      707
      53      413
     112      491
      Name: Rating, Length: 320, dtype: int64)
```

```
[111]: model = Ridge(alpha=.8)
training_set = model.fit(x_train, y_train)
training_set.coef_, training_set.intercept_

training_test = model.predict(x_train)
acc_score_train = accuracy_score(y_train, training_test.astype(int))

print(acc_score_train,mean_squared_error(y_train, training_test.astype(int)))
```

```
0.04375 103.609375
```

```
[103]: # test the model to the testing data.
prediction = model.predict(x_test)
np.array(y_test),prediction.astype(int)

# testing error and metrics.
acc_score = accuracy_score(y_test,prediction.astype(int))
```

```
print(acc_score)

mean_squared_error(y_test, prediction.astype(int))
```

0.025

[103]: 98.475

## 5 Question 2

Choose a classification dataset (not the adult.data set, The UCI repository has many datasets as well as Kaggle), perform test/train split and create a classification model (your choice but DecisionTree is fine). Calculate + Accuracy + Confusion Matrix + Classification Report

```
[126]: # Clean up data.
cancer_data.isin(['?']).any()
index_name = cancer_data[cancer_data['bare_nucleoli'] == '?'].index
cancer_data.drop(index_name, inplace=True)
```

```
[127]: x_train, x_test, y_train, y_test = train_test_split(cancer_data.
↳ drop(['cell_class'], axis=1), cancer_data.cell_class, test_size=.20)
```

```
[128]: # decision tree classification of cancer cells.
model = DecisionTreeClassifier(criterion='entropy')
model.fit(x_train, y_train)
```

```
[128]: DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='entropy',
                             max_depth=None, max_features=None, max_leaf_nodes=None,
                             min_impurity_decrease=0.0, min_impurity_split=None,
                             min_samples_leaf=1, min_samples_split=2,
                             min_weight_fraction_leaf=0.0, presort='deprecated',
                             random_state=None, splitter='best')
```

```
[132]: test_predictions = model.predict(x_test)
acc = accuracy_score(y_test, test_predictions)
conf_matrx = confusion_matrix(y_test, test_predictions)
report = classification_report(y_test, test_predictions)
print(acc)
print(conf_matrx)
print(report)
```

0.9562043795620438

```
[[89  4]
```

```
 [ 2 42]]
```

	precision	recall	f1-score	support
2	0.98	0.96	0.97	93

	4	0.91	0.95	0.93	44
accuracy				0.96	137
macro avg		0.95	0.96	0.95	137
weighted avg		0.96	0.96	0.96	137

## 6 Bonus

(Bonus) See if you can improve the classification model's performance with any tricks you can think of (modify features, remove features, polynomial features)

```
[156]: # modify the alpha value to see an increase in score.
credit.Gender.replace('Male', 1)
credit.Gender.replace('Female', 0)
credit.Student.replace('Yes', 1)
credit.Student.replace('No', 0)
credit.Married.replace('Yes', 1)
credit.Married.replace('No', 0)

numerical_cats = credit.select_dtypes(include=['int64', 'float64'])
x_train, x_test, y_train, y_test = train_test_split(numerical_cats.
    ↳ drop(['Rating'], axis=1), numerical_cats.Rating, test_size=.20)
x_train, y_train

model = Ridge(alpha=.35)
training_set = model.fit(x_train, y_train)
training_set.coef_, training_set.intercept_

training_test = model.predict(x_train)
acc_score_train = accuracy_score(y_train, training_test.astype(int))

print(acc_score_train, mean_squared_error(y_train, training_test.astype(int)))

# test the model to the testing data.
prediction = model.predict(x_test)
np.array(y_test), prediction.astype(int)

# testing error and metrics.
acc_score = accuracy_score(y_test, prediction.astype(int))
print(acc_score)

mean_squared_error(y_test, prediction.astype(int))
```

```
0.05 103.775
0.0125
```

[156]: 98.125

[ ]: