

March 31, 2020

1 9 Clustering

1.1 1. DBSCAN

Using DBSCAN iterate (for-loop) through different values of `min_samples` (1 to 10) and `epsilon` (.05 to .5, in steps of .01) to find clusters in the road-data used in the Lesson and calculate the Silhouette Coeff for `min_samples` and `epsilon`. Plot *one* line plot with the multiple lines generated from the `min_samples` and `epsilon` values. Use a 2D array to store the SilCoeff values, one dimension represents `min_samples`, the other represents `epsilon`.

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn
from mpl_toolkits.mplot3d import Axes3D
from sklearn.cluster import DBSCAN
from sklearn.cluster import KMeans
from sklearn import metrics
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import OneHotEncoder
from sklearn.preprocessing import LabelEncoder
from sklearn.decomposition import PCA
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis

[2]: # import necessary file and drop osm that won't be used for clustering.
road_df = pd.read_csv('../data/3D_spatial_network.txt.gz', header=None,
    ↪names=['osm', 'lat', 'lon', 'alt'])
road_df = road_df.drop(['osm'], axis=1).sample(2000)

# describe dataframe.
print(road_df.describe())
print()

# check for missing values.
missing_data = road_df.isna().sum()
print(missing_data)
print()
```

```
# check column dtypes annd verify correct dtypes.
print(road_df.dtypes)
```

```

           lat           lon           alt
count  2000.000000  2000.000000  2000.000000
mean     9.738332    57.091143    21.906253
std      0.631110     0.289266    18.222317
min      8.146658    56.583746    -1.008307
25%      9.333164    56.859586     6.962991
50%      9.889718    57.047835    17.788992
75%     10.185284    57.318931    31.066453
max     11.196181    57.739892    99.483528
```

```
lat      0
lon      0
alt      0
dtype: int64
```

```
lat      float64
lon      float64
alt      float64
dtype: object
```

```
[3]: # scale all values for clustering.
scaler = StandardScaler()
scaled_values = scaler.fit_transform(road_df)
scaled_df = pd.DataFrame(scaled_values)
scaled_df.columns = ['lat', 'lon', 'alt']
print(scaled_df.head())
```

```

           lat           lon           alt
0  0.194917  1.134797 -0.020862
1  1.218590  1.168535 -0.386648
2  0.343548  1.596188  0.569756
3 -2.322370 -1.729242 -1.151099
4  0.733573 -0.392915 -1.047892
```

```
[4]: eps = .05
generated_values = []
s_scores = []
packed_list = []

while eps < 0.5:
    for i in range(1,11):

        # create dbscan object and cluster samples together.
        dbscan = DBSCAN(eps=eps, min_samples=i)
```

```

scaled_df['cluster'] = dbscan.fit_predict(scaled_df[['lat', 'lon', 'alt',
→ 'alt']])

# append generated_values with each df to plot later.
generated_values.append(scaled_df)
#print(scaled_df.describe())
sil_scores = metrics.silhouette_score(scaled_df[['lon', 'lat', 'alt']],
→ scaled_df.cluster)
s_scores.append(sil_scores)
packed_list.append([i,eps,sil_scores])

#update eps score for the loop to move onto the next iteration.
eps+=.01

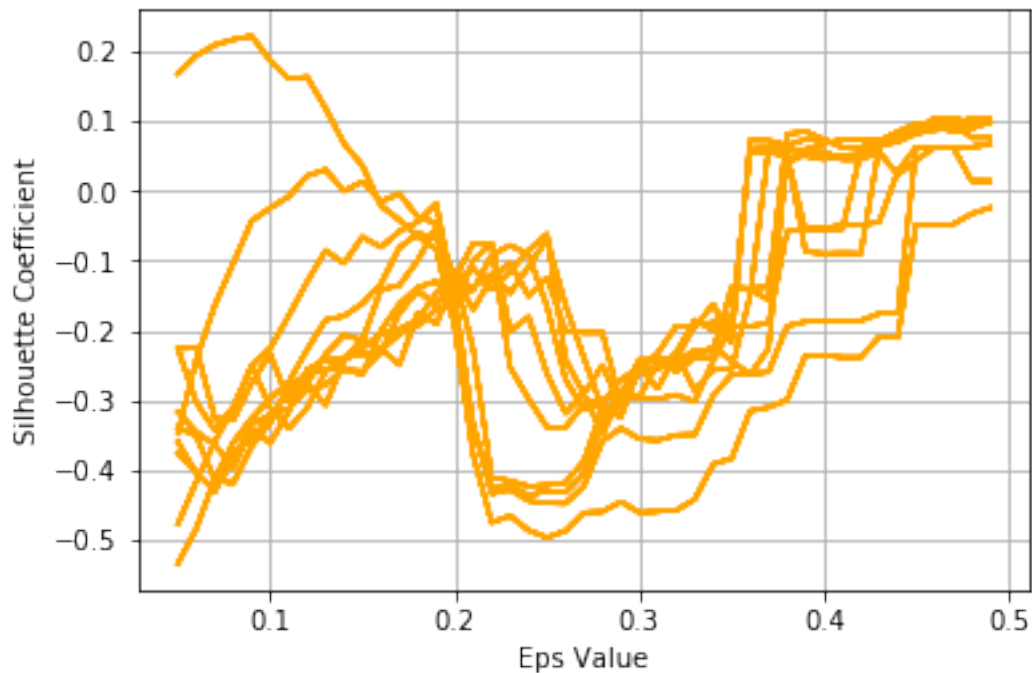
scores_array = np.array(packed_list)

```

```

[5]: y = scores_array[:,2].reshape(-1,10).T
x = scores_array[:,1].reshape(-1,10).T
for row in y:
    for num in x:
        plt.plot(num, row, c = 'orange')
        plt.xlabel('Eps Value')
        plt.ylabel('Silhouette Coefficient')
        plt.grid(True)

```



1.2 2. Clustering your own data

Using your own data, find relevant clusters/groups within your data. If your data is labeled already, with a class that you are attempting to predict, be sure to not use it in fitting/training/predicting.

You may use the labels to compare with predictions to show how well the clustering performed using one of the clustering metrics (<http://scikit-learn.org/stable/modules/clustering.html#clustering-performance-evaluation>).

If you don't have labels, use the silhouette coefficient to show performance. Find the optimal fit for your data but you don't need to be as exhaustive as above.

Additionally, show the clusters in 2D and 3D plots.

For bonus, try using PCA first to condense your data from N columns to less than N.

Two items are expected: - Metric Evaluation Plot - Plots of the clustered data

2 Outline for Answers

2.1 1. Cleaning data

2.2 2. KMeans Clustering w/ Graphing & Metric Eval

2.3 3. DBSCAN Clustering w/Graphing & Metric Eval

2.4 4. PCA

2.5 Data Cleaning

```
[6]: beers_df = pd.read_csv('../data/beers.csv', na_values=['', ' ', 'nan'])

# inspect the data.
beers_df = beers_df.drop(['Unnamed: 0'], axis = 1)
print(beers_df.head(), '\n')
print(beers_df.describe(), '\n')
print(beers_df.isna().sum(), '\n')
print(beers_df.dtypes)
```

	abv	ibu	id	name	style \
0	0.050	NaN	1436	Pub Beer	American Pale Lager
1	0.066	NaN	2265	Devil's Cup	American Pale Ale (APA)
2	0.071	NaN	2264	Rise of the Phoenix	American IPA
3	0.090	NaN	2263	Sinister	American Double / Imperial IPA
4	0.075	NaN	2262	Sex and Candy	American IPA

	brewery_id	ounces
0	408	12.0
1	177	12.0
2	177	12.0

```
3      177    12.0
4      177    12.0
```

```

      abv      ibu      id  brewery_id      ounces
count 2348.000000 1405.000000 2410.000000 2410.000000 2410.000000
mean    0.059773   42.713167 1431.113278   231.749793   13.592241
std     0.013542   25.954066   752.459975   157.685604    2.352204
min     0.001000    4.000000    1.000000    0.000000    8.400000
25%     0.050000   21.000000   808.250000    93.000000   12.000000
50%     0.056000   35.000000  1453.500000   205.000000   12.000000
75%     0.067000   64.000000  2075.750000   366.000000   16.000000
max     0.128000  138.000000  2692.000000   557.000000   32.000000

```

```

abv      62
ibu     1005
id         0
name         0
style        5
brewery_id   0
ounces       0
dtype: int64

```

```

abv      float64
ibu      float64
id        int64
name      object
style     object
brewery_id  int64
ounces    float64
dtype: object

```

```
[7]: # fill in missing values and drop the rest that aren't filled to improve model.
#beers_df['ibu'] = beers_df['ibu'].replace('nan', np.nan, inplace = True)
beers_df['ibu'] = beers_df['ibu'].fillna(beers_df.groupby('style')['ibu'].
    ↳transform('mean'))
beers_df['abv'] = beers_df['abv'].fillna(beers_df.groupby('style')['abv'].
    ↳transform('mean'))
beers_df = beers_df.dropna(axis=0)

beers_df.isna().sum(), beers_df.describe()
```

```
[7]: (abv      0
      ibu      0
      id      0
      name     0
      style    0
      brewery_id 0)
```

```

ounces      0
dtype: int64,
      abv      ibu      id      brewery_id      ounces
count  2353.000000  2353.000000  2353.000000  2353.000000  2353.000000
mean    0.059731    40.900492  1426.521037    230.332342    13.564853
std     0.013443    23.851880   756.589326   157.789185     2.316091
min     0.027000     4.000000     1.000000     0.000000     8.400000
25%     0.050000    22.000000   796.000000    91.000000    12.000000
50%     0.056000    34.125000  1451.000000   204.000000    12.000000
75%     0.067000    60.000000  2075.000000   365.000000    16.000000
max     0.128000   138.000000  2692.000000   557.000000    32.000000

```

```

[8]: # convert categorical data to numerical for later confirmation.
enc = OneHotEncoder()
label_encoder = LabelEncoder()
beer_style = beers_df['style']
beers_df['target_styles'] = label_encoder.fit_transform(beer_style)
beers_df.head()

```

```

[8]:      abv      ibu      id      name \
0  0.050  26.750000  1436      Pub Beer
1  0.066  44.941176  2265      Devil's Cup
2  0.071  67.634551  2264  Rise of the Phoenix
3  0.090  93.320000  2263      Sinister
4  0.075  67.634551  2262      Sex and Candy

      style      brewery_id      ounces      target_styles
0      American Pale Lager      408      12.0      16
1      American Pale Ale (APA)      177      12.0      15
2      American IPA      177      12.0      13
3  American Double / Imperial IPA      177      12.0      10
4      American IPA      177      12.0      13

```

```

[9]: # extract numerical data for clustering
numerical_df = beers_df.select_dtypes(['int64', 'float64']).copy()
numerical_df.head()

```

```

[9]:      abv      ibu      id      brewery_id      ounces      target_styles
0  0.050  26.750000  1436      408      12.0      16
1  0.066  44.941176  2265      177      12.0      15
2  0.071  67.634551  2264      177      12.0      13
3  0.090  93.320000  2263      177      12.0      10
4  0.075  67.634551  2262      177      12.0      13

```

```

[10]: # get unique count for each of beer to check with results.
unique_beer_type = len(numerical_df.target_styles.unique())
print(f'Unique types of beers: {unique_beer_type}')

```

Unique types of beers: 90

```
[11]: # create features df.
features = numerical_df.drop(['target_styles', 'brewery_id', 'id'], axis = 1).
      ↪ copy()
print(features)
# create 2d np array for classification features.
beer_labels = numerical_df.target_styles

# create 1d np array of label converted style of beer.
```

	abv	ibu	ounces
0	0.050	26.750000	12.0
1	0.066	44.941176	12.0
2	0.071	67.634551	12.0
3	0.090	93.320000	12.0
4	0.075	67.634551	12.0
...
2405	0.067	45.000000	12.0
2406	0.052	36.298701	12.0
2407	0.055	31.000000	12.0
2408	0.055	40.000000	12.0
2409	0.052	36.298701	12.0

[2353 rows x 3 columns]

2.6 KMeans Clustering

```
[12]: # get kmeans clusters.
km = KMeans()
km.fit(features)
beer_clusters = km.predict(features)
numerical_df['cluster'] = beer_clusters.copy()
```

```
[13]: counts = np.bincount(numerical_df.cluster)
ii = np.nonzero(counts)[0]
np.vstack((ii, counts[ii])).T
```

```
[13]: array([[ 0, 612],
           [ 1, 345],
           [ 2, 177],
           [ 3,  58],
           [ 4, 443],
           [ 5, 301],
           [ 6, 284],
           [ 7, 133]])
```

2.7 KMeans Metric Evaluation

```
[14]: from sklearn import metrics
      metrics.homogeneity_score(beer_labels, beer_clusters)
```

```
[14]: 0.27540498871449065
```


2.8 KMeans Clustering and Graphing in 2D and 3D

```
[15]: fig = plt.figure()
      plt.clf()
      ax = Axes3D(fig, rect=[0, 0, .95, 1], elev=18, azimuth=150)

      plt.cla()

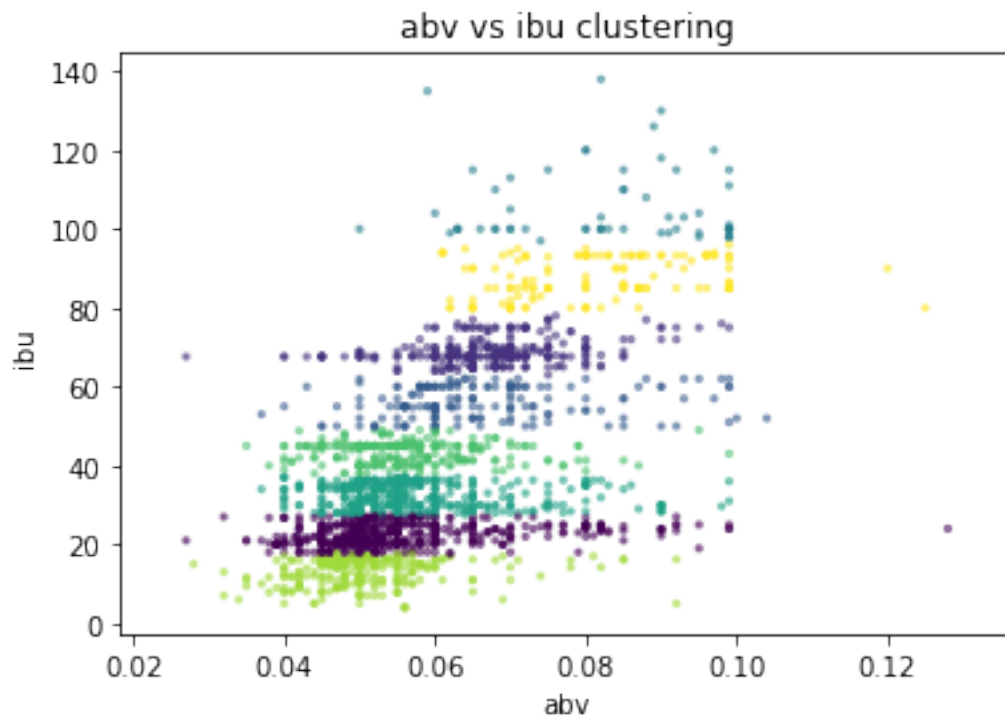
      ax.scatter(numerical_df['abv'],
                  numerical_df['ibu'],
                  numerical_df['ounces'],
                  c=numerical_df.cluster,
                  s=5)

      ax.set_xlabel('abv')
      ax.set_ylabel('ibu')
      ax.set_zlabel('ounces')
      plt.show()
```


```
[16]: plt.scatter(numerical_df['abv'],  
                  numerical_df['ibu'],  
                  s=5,  
                  c=numerical_df.cluster,  
                  alpha=0.5)  
plt.title("abv vs ibu clustering")  
plt.xlabel("abv")  
plt.ylabel("ibu")
```

```
[16]: Text(0,0.5,'ibu')
```




```
[17]: plt.scatter(numerical_df['abv'],  
                 numerical_df['ounces'],  
                 s=5,  
                 c=numerical_df.cluster,  
                 alpha=0.5)  
plt.title("abv vs ounces clustering")  
plt.xlabel("abv")  
plt.ylabel("ounces")
```

```
[17]: Text(0,0.5,'ounces')
```



```
[18]: plt.scatter(numerical_df['ibu'],  
                 numerical_df['ounces'],  
                 s=5,  
                 c=numerical_df.cluster,  
                 alpha=0.5)  
plt.title("ibu vs ounces clustering")  
plt.xlabel("ibu")  
plt.ylabel("ounces")
```

```
[18]: Text(0,0.5,'ounces')
```



```
[20]: # normalize the data
# scale all values for clustering.
scaler = StandardScaler()
scaled_values = scaler.fit_transform(features)
scaled_beer_features = pd.DataFrame(scaled_values)
scaled_beer_features.columns = ['abv', 'ibu', 'ounces']
scaled_beer_features = scaled_beer_features.rename(columns={0: 'abv', 1: 'ibu', 2:
    → 'ounces'})
print(scaled_beer_features.head(20))
```

	abv	ibu	ounces
0	-0.723979	-0.593391	-0.675788
1	0.466453	0.169443	-0.675788
2	0.838463	1.121075	-0.675788
3	2.252101	2.198177	-0.675788
4	1.136071	1.121075	-0.675788
5	1.284875	-0.537161	-0.675788
6	-1.095989	0.169443	-0.675788
7	0.392051	-0.376462	-0.675788
8	-0.351969	0.169443	-0.675788
9	1.954493	2.198177	-0.675788
10	0.912865	-0.524564	-0.675788
11	0.987267	-0.524564	-0.675788
12	0.689659	-0.524564	-0.675788

```
13  1.880091  0.675122 -0.675788
14  0.094443  0.800925 -0.675788
15  1.656885  0.549319 -0.675788
16  1.656885 -0.729678 -0.675788
17  2.921719  2.142824 -2.230461
18  1.433679  0.171910 -0.675788
19  1.433679  0.017277 -0.675788
```

/Users/Carancho/miniconda3/envs/ml_class_work/lib/python3.6/site-packages/ipykernel_launcher.py:6: UserWarning: Pandas doesn't allow columns to be created via a new attribute name - see <https://pandas.pydata.org/pandas-docs/stable/indexing.html#attribute-access>


```
[21]: # get kmeans clusters.
km.fit(scaled_beer_features)
scaled_clusters = km.predict(scaled_beer_features)
scaled_beer_features['kmeans_cluster'] = scaled_clusters.copy()

fig = plt.figure()
plt.clf()
ax = Axes3D(fig, rect=[0, 0, .95, 1], elev=18, azimuth=150)

plt.cla()

ax.scatter(scaled_beer_features['abv'],
           scaled_beer_features['ibu'],
           scaled_beer_features['ounces'],
           c=scaled_beer_features.kmeans_cluster,
           s=5)

ax.set_xlabel('abv')
ax.set_ylabel('ibu')
ax.set_zlabel('ounces')
plt.show()
```




```
[22]: metrics.homogeneity_score(beer_labels, scaled_clusters)
```

```
[22]: 0.21773268370467797
```

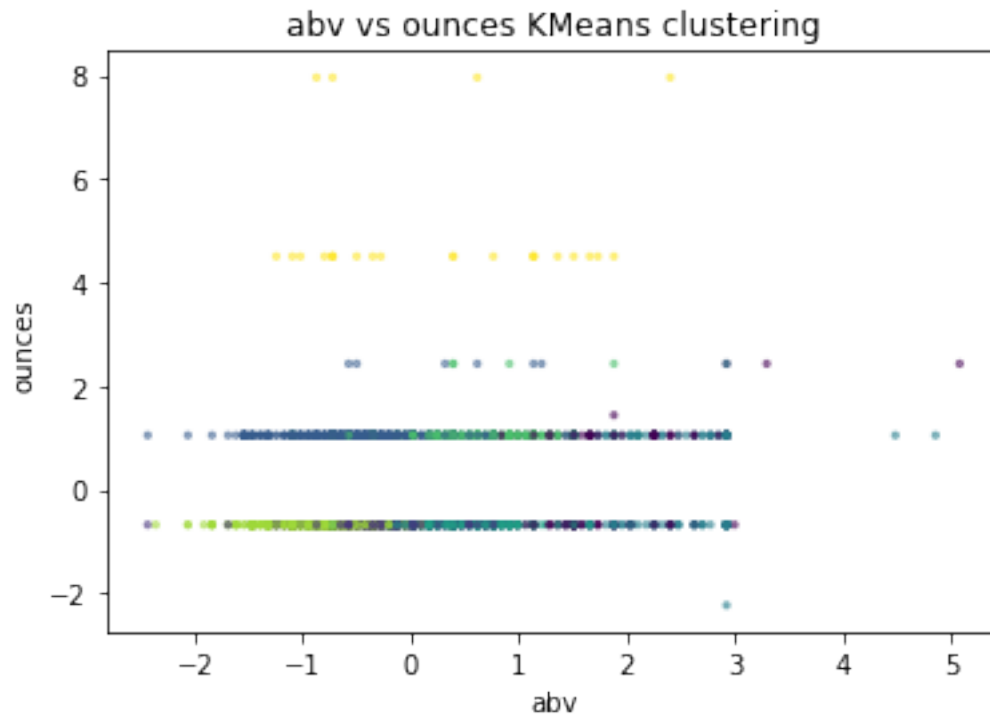
```
[23]: plt.scatter(scaled_beer_features['abv'],
                  scaled_beer_features['ibu'],
                  s=5,
                  c=scaled_beer_features.kmeans_cluster,
                  alpha=0.5)
plt.title("abv vs ibu KMeans clustering")
plt.xlabel("abv")
plt.ylabel("ibu")
```

```
[23]: Text(0,0.5,'ibu')
```




```
[24]: plt.scatter(scaled_beer_features['abv'],
                  scaled_beer_features['ounces'],
                  s=5,
                  c=scaled_beer_features.kmeans_cluster,
                  alpha=0.5)
plt.title("abv vs ounces KMeans clustering")
plt.xlabel("abv")
plt.ylabel("ounces")
```

```
[24]: Text(0,0.5,'ounces')
```



```
[25]: plt.scatter(scaled_beer_features['ibu'],
                  scaled_beer_features['ounces'],
                  s=5,
                  c=scaled_beer_features.kmeans_cluster,
                  alpha=0.5)
plt.title("ibu vs ounces KMeans clustering")
plt.xlabel("ibu")
plt.ylabel("ounces")
```

```
[25]: Text(0,0.5,'ounces')
```

2.9 DBSCAN Clustering and Graphing in 2D and 3D

```
[26]: # Cluster using DBSCAN.
dbscan = DBSCAN(eps=.5)
dbscan_cluster = dbscan.fit_predict(scaled_beer_features[['abv', 'ibu', '
↪ 'ounces']])
scaled_beer_features['dbscan_cluster'] = dbscan_cluster.copy()
print(scaled_beer_features.head())
```

	abv	ibu	ounces	kmeans_cluster	dbscan_cluster
0	-0.723979	-0.593391	-0.675788	6	0
1	0.466453	0.169443	-0.675788	1	0
2	0.838463	1.121075	-0.675788	4	0
3	2.252101	2.198177	-0.675788	3	0
4	1.136071	1.121075	-0.675788	4	0

```
[27]: metrics.homogeneity_score(beer_labels, dbscan_cluster)
```


```
[27]: 0.020337304064643697
```

```
[28]: fig = plt.figure()
plt.clf()
ax = Axes3D(fig, rect=[0, 0, .95, 1], elev=18, azimuth=150)
```

```
plt.cla()


ax.scatter(scaled_beer_features['abv'],
          scaled_beer_features['ibu'],
          scaled_beer_features['ounces'],
          c=scaled_beer_features.dbscan_cluster,
          s=5)

ax.set_xlabel('abv')
ax.set_ylabel('ibu')
ax.set_zlabel('ounces')
plt.show()
```



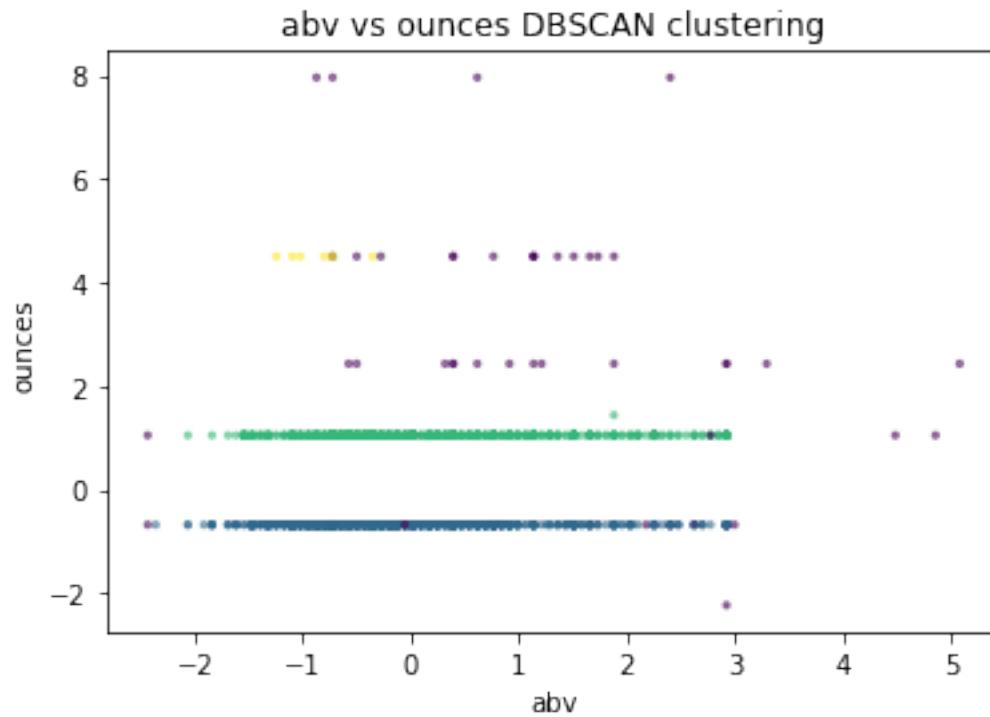
```
[29]: plt.scatter(scaled_beer_features['abv'],
                  scaled_beer_features['ibu'],
                  s=5,
                  c=scaled_beer_features.dbscan_cluster,
                  alpha=0.5)
plt.title("abv vs ibu DBSCAN clustering")
plt.xlabel("abv")
plt.ylabel("ibu")
```

```
[29]: Text(0,0.5,'ibu')
```




```
[30]: plt.scatter(scaled_beer_features['abv'],
                  scaled_beer_features['ounces'],
                  s=5,
                  c=scaled_beer_features.dbscan_cluster,
                  alpha=0.5)
plt.title("abv vs ounces DBSCAN clustering")
plt.xlabel("abv")
plt.ylabel("ounces")
```

```
[30]: Text(0,0.5,'ounces')
```



```
[31]: plt.scatter(scaled_beer_features['ibu'],
                  scaled_beer_features['ounces'],
                  s=5,
                  c=scaled_beer_features.dbscan_cluster,
                  alpha=0.5)
plt.title("ibu vs ounces DBSCAN clustering")
plt.xlabel("ibu")
plt.ylabel("ounces")
```

```
[31]: Text(0,0.5,'ounces')
```



2.10 PCA

```
[32]: pca = PCA(n_components=2)
X_r = pca.fit(scaled_beer_features).transform(scaled_beer_features)
pca_df = pd.DataFrame(X_r)
pca_df['style'] = beer_style
lda = LinearDiscriminantAnalysis(n_components=2)
X_r2 = lda.fit(scaled_beer_features, target_styles).
↳transform(scaled_beer_features)
lda_df = pd.DataFrame(X_r2)
```

```
↳
-----
NameError                                Traceback (most recent call↳
↳last)

<ipython-input-32-4a2940a2e706> in <module>
      4 pca_df['style'] = beer_style
      5 lda = LinearDiscriminantAnalysis(n_components=2)
----> 6 X_r2 = lda.fit(scaled_beer_features, target_styles).
↳transform(scaled_beer_features)
```

```
7 lda_df = pd.DataFrame(X_r2)
```

NameError: name 'target_styles' is not defined

```
[33]: print('explained variance ratio (first two components): %s'
      % str(pca.explained_variance_ratio_))
```

explained variance ratio (first two components): [0.58329667 0.20850056]

```
[34]: pca_df.head()
```


```
[34]:      0      1      style
0 -2.818627 -0.686346      American Pale Lager
1  2.314626 -0.121596      American Pale Ale (APA)
2 -0.544530  1.223370      American IPA
3  0.699529  2.743023  American Double / Imperial IPA
4 -0.497882  1.417579      American IPA
```

```
[35]: len(beer_labels.unique())
```

```
[35]: 90
```

```
[36]: plt.scatter(pca_df.iloc[:,0], pca_df.iloc[:,1], c=np.random.rand(2353), label =_
      ↪beer_labels)
```

```
[36]: <matplotlib.collections.PathCollection at 0x1a1ecd8588>
```



2.11 Note

You may use any for both parts 1 and 2, I only recommend using the data I used in the Lesson for part 1. I've included several new datasets in the `data/` folder, such as `beers.csv`, `snow_tweets.csv`, `data/USCensus1990.data.txt.gz`. You do not need to unzip or ungzip any data files. Pandas can open these files on its own.