

April 11, 2020

## 1 Assignment 12 - Neural Networks image recognition

Use both MLNN and the ConvNet to solve the following problem.

1. Add random noise (i.e. `np.random.normal`) to the images in training and testing. Make sure each image gets a different noise feature added to it. Inspect by printing out an image.
2. Compare the loss/accuracy (train, val) after N epochs for both MLNN and ConvNet with and without noise.
3. Vary the amount of noise (multiply `np.random.normal` by a factor) and keep track of the accuracy and loss (for training and validation) and plot these results.

## 2 Neural Networks - Image Recognition

```
[2]: import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.optimizers import RMSprop
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras import backend
import numpy as np
from skimage.util import random_noise
import matplotlib.pyplot as plt
%matplotlib inline
```

### 2.1 Multi Layer Neural Network

Trains a simple deep NN on the MNIST dataset. Gets to 98.40% test accuracy after 20 epochs (there is *a lot* of margin for parameter tuning).

```
[13]: # the data, shuffled and split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()


x_train = x_train.reshape(60000, 784)
x_test = x_test.reshape(10000, 784)
x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
```

```
x_test /= 255
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')
print(x_train)
```

```
60000 train samples
10000 test samples
[[0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]
 ...
 [0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]]
```

```
[14]: plt.imshow(x_train[2].reshape(28,28))
```

```
[14]: <matplotlib.image.AxesImage at 0x63739cba8>
```



```
[15]: initial_mlnn_loss = []
```

```
[16]: batch_size = 128
      num_classes = 10
      epochs = 20
```

```

# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

model = Sequential()
model.add(Dense(512, activation='relu', input_shape=(784,)))
model.add(Dropout(0.2))
model.add(Dense(512, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(10, activation='softmax'))

model.summary()

model.compile(loss='categorical_crossentropy',
              optimizer=RMSprop(),
              metrics=['accuracy'])

history = model.fit(x_train, y_train,
                    batch_size=batch_size,
                    epochs=epochs,
                    verbose=1,
                    validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
initial_mlnn_loss.append(model.evaluate(x_test, y_test, verbose=0))
print('Test loss:', score[0])
print('Test accuracy:', score[1])

```

Model: "sequential\_3"

Layer (type)	Output Shape	Param #
dense_7 (Dense)	(None, 512)	401920
dropout_5 (Dropout)	(None, 512)	0
dense_8 (Dense)	(None, 512)	262656
dropout_6 (Dropout)	(None, 512)	0
dense_9 (Dense)	(None, 10)	5130

Total params: 669,706  
 Trainable params: 669,706  
 Non-trainable params: 0

Train on 60000 samples, validate on 10000 samples

Epoch 1/20  
60000/60000 [=====] - 3s 43us/step - loss: 0.2449 - accuracy: 0.9245 - val\_loss: 0.1072 - val\_accuracy: 0.9669

Epoch 2/20  
60000/60000 [=====] - 3s 42us/step - loss: 0.1039 - accuracy: 0.9683 - val\_loss: 0.0888 - val\_accuracy: 0.9725

Epoch 3/20  
60000/60000 [=====] - 3s 42us/step - loss: 0.0748 - accuracy: 0.9781 - val\_loss: 0.0774 - val\_accuracy: 0.9786

Epoch 4/20  
60000/60000 [=====] - 3s 43us/step - loss: 0.0595 - accuracy: 0.9820 - val\_loss: 0.0987 - val\_accuracy: 0.9720

Epoch 5/20  
60000/60000 [=====] - 3s 45us/step - loss: 0.0489 - accuracy: 0.9855 - val\_loss: 0.0780 - val\_accuracy: 0.9795

Epoch 6/20  
60000/60000 [=====] - 3s 43us/step - loss: 0.0439 - accuracy: 0.9869 - val\_loss: 0.0935 - val\_accuracy: 0.9777

Epoch 7/20  
60000/60000 [=====] - 2s 42us/step - loss: 0.0380 - accuracy: 0.9883 - val\_loss: 0.0761 - val\_accuracy: 0.9827

Epoch 8/20  
60000/60000 [=====] - 2s 42us/step - loss: 0.0350 - accuracy: 0.9898 - val\_loss: 0.0841 - val\_accuracy: 0.9815

Epoch 9/20  
60000/60000 [=====] - 3s 48us/step - loss: 0.0301 - accuracy: 0.9908 - val\_loss: 0.0770 - val\_accuracy: 0.9842

Epoch 10/20  
60000/60000 [=====] - 3s 55us/step - loss: 0.0287 - accuracy: 0.9924 - val\_loss: 0.0910 - val\_accuracy: 0.9832

Epoch 11/20  
60000/60000 [=====] - 3s 44us/step - loss: 0.0281 - accuracy: 0.9915 - val\_loss: 0.0950 - val\_accuracy: 0.9824

Epoch 12/20  
60000/60000 [=====] - 3s 42us/step - loss: 0.0243 - accuracy: 0.9933 - val\_loss: 0.0963 - val\_accuracy: 0.9815

Epoch 13/20  
60000/60000 [=====] - 3s 44us/step - loss: 0.0243 - accuracy: 0.9931 - val\_loss: 0.1059 - val\_accuracy: 0.9811

Epoch 14/20  
60000/60000 [=====] - 3s 45us/step - loss: 0.0225 - accuracy: 0.9937 - val\_loss: 0.1000 - val\_accuracy: 0.9834

Epoch 15/20  
60000/60000 [=====] - 3s 47us/step - loss: 0.0197 - accuracy: 0.9942 - val\_loss: 0.0985 - val\_accuracy: 0.9846

Epoch 16/20  
60000/60000 [=====] - 3s 48us/step - loss: 0.0213 - accuracy: 0.9945 - val\_loss: 0.1201 - val\_accuracy: 0.9825

```

Epoch 17/20
60000/60000 [=====] - 3s 43us/step - loss: 0.0213 -
accuracy: 0.9944 - val_loss: 0.1118 - val_accuracy: 0.9821
Epoch 18/20
60000/60000 [=====] - 3s 43us/step - loss: 0.0205 -
accuracy: 0.9944 - val_loss: 0.1155 - val_accuracy: 0.9832
Epoch 19/20
60000/60000 [=====] - 3s 43us/step - loss: 0.0181 -
accuracy: 0.9953 - val_loss: 0.1120 - val_accuracy: 0.9831
Epoch 20/20
60000/60000 [=====] - 2s 41us/step - loss: 0.0181 -
accuracy: 0.9953 - val_loss: 0.1189 - val_accuracy: 0.9826
Test loss: 0.1189450991806239
Test accuracy: 0.9825999736785889

```

```
[10]: initial_mlnn_loss
```

```
[10]: [[0.10650663499242005, 0.983299970626831]]
```

## 2.2 Adding noise to the data.

```

[9]: # the data, shuffled and split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()

x_train = x_train.reshape(60000, 784)
x_train = random_noise(x_train)
x_test = x_test.reshape(10000, 784)
x_test = random_noise(x_test)
x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')
#x_train[0] 60000 x_train[1] 784

```


```

60000 train samples
10000 test samples

```

```
[10]: plt.imshow(x_train[2].reshape(28,28))
```

```
[10]: <matplotlib.image.AxesImage at 0x63405cc18>
```



```
[11]: noisy_mlnn_loss = []
```

```
[12]: batch_size = 128
num_classes = 10
epochs = 20

# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

model = Sequential()
model.add(Dense(512, activation='relu', input_shape=(784,)))
model.add(Dropout(0.2))
model.add(Dense(512, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(10, activation='softmax'))

model.summary()

model.compile(loss='categorical_crossentropy',
              optimizer=RMSprop(),
              metrics=['accuracy'])

history = model.fit(x_train, y_train,
```

```

        batch_size=batch_size,
        epochs=epochs,
        verbose=1,
        validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
noisy_mlnn_loss.append(model.evaluate(x_test, y_test, verbose=0))
print('Test loss:', score[0])
print('Test accuracy:', score[1])

```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
dense_4 (Dense)	(None, 512)	401920
dropout_3 (Dropout)	(None, 512)	0
dense_5 (Dense)	(None, 512)	262656
dropout_4 (Dropout)	(None, 512)	0
dense_6 (Dense)	(None, 10)	5130

Total params: 669,706

Trainable params: 669,706

Non-trainable params: 0

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 2s 41us/step - loss: 0.9061 - accuracy: 0.7138 - val\_loss: 0.4470 - val\_accuracy: 0.8601

Epoch 2/20

60000/60000 [=====] - 2s 38us/step - loss: 0.4005 - accuracy: 0.8796 - val\_loss: 0.3232 - val\_accuracy: 0.9084

Epoch 3/20

60000/60000 [=====] - 2s 38us/step - loss: 0.3067 - accuracy: 0.9098 - val\_loss: 0.2485 - val\_accuracy: 0.9235

Epoch 4/20

60000/60000 [=====] - 3s 42us/step - loss: 0.2444 - accuracy: 0.9267 - val\_loss: 0.2215 - val\_accuracy: 0.9324

Epoch 5/20

60000/60000 [=====] - 3s 43us/step - loss: 0.2018 - accuracy: 0.9398 - val\_loss: 0.1726 - val\_accuracy: 0.9478

Epoch 6/20

60000/60000 [=====] - 2s 41us/step - loss: 0.1695 - accuracy: 0.9495 - val\_loss: 0.1597 - val\_accuracy: 0.9490

Epoch 7/20

60000/60000 [=====] - 3s 42us/step - loss: 0.1467 -  
accuracy: 0.9556 - val\_loss: 0.1440 - val\_accuracy: 0.9554  
Epoch 8/20  
60000/60000 [=====] - 3s 43us/step - loss: 0.1287 -  
accuracy: 0.9614 - val\_loss: 0.1300 - val\_accuracy: 0.9593  
Epoch 9/20  
60000/60000 [=====] - 3s 42us/step - loss: 0.1144 -  
accuracy: 0.9656 - val\_loss: 0.1129 - val\_accuracy: 0.9657  
Epoch 10/20  
60000/60000 [=====] - 3s 43us/step - loss: 0.1028 -  
accuracy: 0.9690 - val\_loss: 0.1168 - val\_accuracy: 0.9650  
Epoch 11/20  
60000/60000 [=====] - 2s 42us/step - loss: 0.0942 -  
accuracy: 0.9707 - val\_loss: 0.1052 - val\_accuracy: 0.9664  
Epoch 12/20  
60000/60000 [=====] - 3s 42us/step - loss: 0.0851 -  
accuracy: 0.9740 - val\_loss: 0.0959 - val\_accuracy: 0.9702  
Epoch 13/20  
60000/60000 [=====] - 3s 43us/step - loss: 0.0768 -  
accuracy: 0.9768 - val\_loss: 0.0939 - val\_accuracy: 0.9719  
Epoch 14/20  
60000/60000 [=====] - 2s 41us/step - loss: 0.0706 -  
accuracy: 0.9785 - val\_loss: 0.0968 - val\_accuracy: 0.9694  
Epoch 15/20  
60000/60000 [=====] - 3s 43us/step - loss: 0.0661 -  
accuracy: 0.9797 - val\_loss: 0.0882 - val\_accuracy: 0.9739  
Epoch 16/20  
60000/60000 [=====] - 3s 42us/step - loss: 0.0616 -  
accuracy: 0.9813 - val\_loss: 0.0856 - val\_accuracy: 0.9740  
Epoch 17/20  
60000/60000 [=====] - 3s 42us/step - loss: 0.0548 -  
accuracy: 0.9832 - val\_loss: 0.0900 - val\_accuracy: 0.9740  
Epoch 18/20  
60000/60000 [=====] - 3s 42us/step - loss: 0.0520 -  
accuracy: 0.9842 - val\_loss: 0.0930 - val\_accuracy: 0.9732  
Epoch 19/20  
60000/60000 [=====] - 3s 43us/step - loss: 0.0486 -  
accuracy: 0.9849 - val\_loss: 0.0916 - val\_accuracy: 0.9748  
Epoch 20/20  
60000/60000 [=====] - 3s 42us/step - loss: 0.0457 -  
accuracy: 0.9862 - val\_loss: 0.0853 - val\_accuracy: 0.9761  
Test loss: 0.08534564933415968  
Test accuracy: 0.9761000275611877



## 2.3 Adding more noise

```
[31]: # the data, shuffled and split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()

x_train = x_train.reshape(60000, 784)
x_train = random_noise(x_train, mode='speckle')
x_test = x_test.reshape(10000, 784)
x_test = random_noise(x_test, mode='speckle')
x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')
#x_train[0] 60000 x_train[1] 784
```

60000 train samples  
10000 test samples

```
[32]: more_noise_loss = []
```

```
[33]: batch_size = 128
num_classes = 10
epochs = 20

# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

model = Sequential()
model.add(Dense(512, activation='relu', input_shape=(784,)))
model.add(Dropout(0.2))
model.add(Dense(512, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(10, activation='softmax'))

model.summary()

model.compile(loss='categorical_crossentropy',
              optimizer=RMSprop(),
              metrics=['accuracy'])

history = model.fit(x_train, y_train,
                    batch_size=batch_size,
                    epochs=epochs,
```

```

        verbose=1,
        validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
more_noise_loss.append(model.evaluate(x_test, y_test, verbose=0))
print('Test loss:', score[0])
print('Test accuracy:', score[1])

```

Model: "sequential\_7"

Layer (type)	Output Shape	Param #
dense_16 (Dense)	(None, 512)	401920
dropout_13 (Dropout)	(None, 512)	0
dense_17 (Dense)	(None, 512)	262656
dropout_14 (Dropout)	(None, 512)	0
dense_18 (Dense)	(None, 10)	5130

Total params: 669,706

Trainable params: 669,706

Non-trainable params: 0

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 3s 45us/step - loss: 0.8786 - accuracy: 0.7274 - val\_loss: 0.4054 - val\_accuracy: 0.8766

Epoch 2/20

60000/60000 [=====] - 2s 39us/step - loss: 0.3737 - accuracy: 0.8879 - val\_loss: 0.2945 - val\_accuracy: 0.9150

Epoch 3/20

60000/60000 [=====] - 2s 39us/step - loss: 0.2832 - accuracy: 0.9154 - val\_loss: 0.2380 - val\_accuracy: 0.9265

Epoch 4/20

60000/60000 [=====] - 2s 41us/step - loss: 0.2257 - accuracy: 0.9323 - val\_loss: 0.1906 - val\_accuracy: 0.9429

Epoch 5/20

60000/60000 [=====] - 3s 44us/step - loss: 0.1873 - accuracy: 0.9426 - val\_loss: 0.1701 - val\_accuracy: 0.9485

Epoch 6/20

60000/60000 [=====] - 3s 43us/step - loss: 0.1591 - accuracy: 0.9521 - val\_loss: 0.1553 - val\_accuracy: 0.9522

Epoch 7/20

60000/60000 [=====] - 3s 42us/step - loss: 0.1406 - accuracy: 0.9578 - val\_loss: 0.1257 - val\_accuracy: 0.9628

```

Epoch 8/20
60000/60000 [=====] - 3s 42us/step - loss: 0.1217 -
accuracy: 0.9634 - val_loss: 0.1173 - val_accuracy: 0.9638
Epoch 9/20
60000/60000 [=====] - 3s 43us/step - loss: 0.1114 -
accuracy: 0.9658 - val_loss: 0.1087 - val_accuracy: 0.9663
Epoch 10/20
60000/60000 [=====] - 3s 43us/step - loss: 0.1017 -
accuracy: 0.9690 - val_loss: 0.1096 - val_accuracy: 0.9668
Epoch 11/20
60000/60000 [=====] - 3s 44us/step - loss: 0.0943 -
accuracy: 0.9718 - val_loss: 0.0944 - val_accuracy: 0.9699
Epoch 12/20
60000/60000 [=====] - 3s 43us/step - loss: 0.0872 -
accuracy: 0.9739 - val_loss: 0.0873 - val_accuracy: 0.9747
Epoch 13/20
60000/60000 [=====] - 3s 42us/step - loss: 0.0800 -
accuracy: 0.9757 - val_loss: 0.0914 - val_accuracy: 0.9736
Epoch 14/20
60000/60000 [=====] - 3s 43us/step - loss: 0.0751 -
accuracy: 0.9765 - val_loss: 0.0875 - val_accuracy: 0.9747
Epoch 15/20
60000/60000 [=====] - 3s 42us/step - loss: 0.0709 -
accuracy: 0.9784 - val_loss: 0.0774 - val_accuracy: 0.9772
Epoch 16/20
60000/60000 [=====] - 3s 44us/step - loss: 0.0666 -
accuracy: 0.9804 - val_loss: 0.0831 - val_accuracy: 0.9763
Epoch 17/20
60000/60000 [=====] - 3s 43us/step - loss: 0.0631 -
accuracy: 0.9803 - val_loss: 0.0825 - val_accuracy: 0.9756
Epoch 18/20
60000/60000 [=====] - 3s 42us/step - loss: 0.0601 -
accuracy: 0.9815 - val_loss: 0.0956 - val_accuracy: 0.9712
Epoch 19/20
60000/60000 [=====] - 3s 43us/step - loss: 0.0568 -
accuracy: 0.9825 - val_loss: 0.0792 - val_accuracy: 0.9780
Epoch 20/20
60000/60000 [=====] - 3s 43us/step - loss: 0.0547 -
accuracy: 0.9834 - val_loss: 0.0823 - val_accuracy: 0.9770
Test loss: 0.08227490533783566
Test accuracy: 0.976999980926514

```

```

[34]: print(noisey_mlnn_loss)
      print(initial_mlnn_loss)
      print(more_noise_loss)

```

```

[[0.08534564933415968, 0.9761000275611877]]
[[0.1189450991806239, 0.9825999736785889]]


```

```
[[0.08227490533783566, 0.9769999980926514]]
```

## 2.4 Graphing loss scores for MLNN

```
[63]: # flatten the lists containing loss scores.  
x1 = np.array(noisey_mlenn_loss).flatten()[0]  
y1 = np.array(noisey_mlenn_loss).flatten()[1]  
  
x2 = np.array(initial_mlenn_loss).flatten()[0]  
y2 = np.array(initial_mlenn_loss).flatten()[1]  
  
x3 = np.array(more_noise_loss).flatten()[0]  
y3 = np.array(more_noise_loss).flatten()[1]  
  
x = [x2,x1,x3]  
y = [y2,y1,y3]  
# graph the data.  
plt.plot(x,y)  
plt.xlabel("Loss Score for MLNN")  
plt.ylabel("Accuracy Score for MLNN")
```

```
[63]: Text(0,0.5,'Accuracy Score for MLNN')
```



## 2.5 Conv Net

Trains a simple convnet on the MNIST dataset. Gets to 99.25% test accuracy after 12 epochs (there is still a lot of margin for parameter tuning).

```
[18]: convnet_loss = []
```

```
[19]: # input image dimensions
img_rows, img_cols = 28, 28

# the data, shuffled and split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()

if backend.image_data_format() == 'channels_first':
    x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
    x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
    input_shape = (1, img_rows, img_cols)
else:
    x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
    x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
    input_shape = (img_rows, img_cols, 1)

x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print('x_train shape:', x_train.shape)
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')
```

```
x_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples
```

```
[20]: batch_size = 128
num_classes = 10
epochs = 12

# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

model = Sequential()
model.add(Conv2D(32, kernel_size=(3, 3),
                 activation='relu',
                 input_shape=input_shape))
model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
```

```

model.add(Dropout(0.25))
model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))

model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Adadelta(),
              metrics=['accuracy'])

model.fit(x_train, y_train,
        batch_size=batch_size,
        epochs=epochs,
        verbose=1,
        validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
convnet_loss.append(model.evaluate(x_test, y_test, verbose=0))
print('Test loss:', score[0])
print('Test accuracy:', score[1])

```

Train on 60000 samples, validate on 10000 samples

Epoch 1/12

60000/60000 [=====] - 62s 1ms/step - loss: 0.2690 - accuracy: 0.9176 - val\_loss: 0.0595 - val\_accuracy: 0.9802

Epoch 2/12

60000/60000 [=====] - 63s 1ms/step - loss: 0.0872 - accuracy: 0.9737 - val\_loss: 0.0429 - val\_accuracy: 0.9854

Epoch 3/12

60000/60000 [=====] - 62s 1ms/step - loss: 0.0662 - accuracy: 0.9796 - val\_loss: 0.0333 - val\_accuracy: 0.9887

Epoch 4/12

60000/60000 [=====] - 61s 1ms/step - loss: 0.0549 - accuracy: 0.9835 - val\_loss: 0.0304 - val\_accuracy: 0.9888

Epoch 5/12

60000/60000 [=====] - 60s 1ms/step - loss: 0.0463 - accuracy: 0.9855 - val\_loss: 0.0331 - val\_accuracy: 0.9882

Epoch 6/12

60000/60000 [=====] - 62s 1ms/step - loss: 0.0420 - accuracy: 0.9873 - val\_loss: 0.0287 - val\_accuracy: 0.9900

Epoch 7/12

60000/60000 [=====] - 61s 1ms/step - loss: 0.0370 - accuracy: 0.9890 - val\_loss: 0.0285 - val\_accuracy: 0.9909

Epoch 8/12

60000/60000 [=====] - 60s 1ms/step - loss: 0.0355 - accuracy: 0.9893 - val\_loss: 0.0291 - val\_accuracy: 0.9911

Epoch 9/12

60000/60000 [=====] - 60s 1ms/step - loss: 0.0337 -

```

accuracy: 0.9895 - val_loss: 0.0331 - val_accuracy: 0.9902
Epoch 10/12
60000/60000 [=====] - 60s 1000us/step - loss: 0.0306 -
accuracy: 0.9904 - val_loss: 0.0245 - val_accuracy: 0.9919
Epoch 11/12
60000/60000 [=====] - 60s 999us/step - loss: 0.0291 -
accuracy: 0.9909 - val_loss: 0.0258 - val_accuracy: 0.9918
Epoch 12/12
60000/60000 [=====] - 60s 1ms/step - loss: 0.0255 -
accuracy: 0.9919 - val_loss: 0.0276 - val_accuracy: 0.9915
Test loss: 0.027638795424243655
Test accuracy: 0.9915000200271606

```

## 2.6 Adding noise to convnet

```
[21]: noise_convnet = []
```

```

[23]: # input image dimensions
img_rows, img_cols = 28, 28

# the data, shuffled and split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()

if backend.image_data_format() == 'channels_first':
    x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
    x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
    input_shape = (1, img_rows, img_cols)
else:
    x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
    x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
    input_shape = (img_rows, img_cols, 1)

x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
x_train = random_noise(x_train)
x_test = random_noise(x_test)

```

```

[24]: batch_size = 128
num_classes = 10
epochs = 12

# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

```

```

model = Sequential()
model.add(Conv2D(32, kernel_size=(3, 3),
                 activation='relu',
                 input_shape=input_shape))
model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))
model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))

model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Adadelta(),
              metrics=['accuracy'])

model.fit(x_train, y_train,
          batch_size=batch_size,
          epochs=epochs,
          verbose=1,
          validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
noise_convnet.append(model.evaluate(x_test, y_test, verbose=0))
print('Test loss:', score[0])
print('Test accuracy:', score[1])

```

Train on 60000 samples, validate on 10000 samples

Epoch 1/12

60000/60000 [=====] - 58s 974us/step - loss: 0.2693 - accuracy: 0.9195 - val\_loss: 0.0592 - val\_accuracy: 0.9820

Epoch 2/12

60000/60000 [=====] - 59s 984us/step - loss: 0.0901 - accuracy: 0.9731 - val\_loss: 0.0402 - val\_accuracy: 0.9861

Epoch 3/12

60000/60000 [=====] - 59s 985us/step - loss: 0.0692 - accuracy: 0.9794 - val\_loss: 0.0371 - val\_accuracy: 0.9881

Epoch 4/12

60000/60000 [=====] - 59s 986us/step - loss: 0.0566 - accuracy: 0.9835 - val\_loss: 0.0308 - val\_accuracy: 0.9895

Epoch 5/12

60000/60000 [=====] - 59s 983us/step - loss: 0.0490 - accuracy: 0.9850 - val\_loss: 0.0327 - val\_accuracy: 0.9884

Epoch 6/12

60000/60000 [=====] - 59s 983us/step - loss: 0.0434 - accuracy: 0.9866 - val\_loss: 0.0316 - val\_accuracy: 0.9906

Epoch 7/12

60000/60000 [=====] - 59s 984us/step - loss: 0.0395 -



```

accuracy: 0.9881 - val_loss: 0.0296 - val_accuracy: 0.9909
Epoch 8/12
60000/60000 [=====] - 59s 983us/step - loss: 0.0356 -
accuracy: 0.9893 - val_loss: 0.0293 - val_accuracy: 0.9909
Epoch 9/12
60000/60000 [=====] - 59s 982us/step - loss: 0.0325 -
accuracy: 0.9901 - val_loss: 0.0291 - val_accuracy: 0.9910
Epoch 10/12
60000/60000 [=====] - 59s 985us/step - loss: 0.0304 -
accuracy: 0.9906 - val_loss: 0.0282 - val_accuracy: 0.9915
Epoch 11/12
60000/60000 [=====] - 59s 984us/step - loss: 0.0280 -
accuracy: 0.9909 - val_loss: 0.0302 - val_accuracy: 0.9913
Epoch 12/12
60000/60000 [=====] - 65s 1ms/step - loss: 0.0259 -
accuracy: 0.9918 - val_loss: 0.0314 - val_accuracy: 0.9905
Test loss: 0.03143596427084858
Test accuracy: 0.9904999732971191

```

## 2.7 Changing the amount of noise added to the training and testing set.

```

[27]: # input image dimensions
img_rows, img_cols = 28, 28

# the data, shuffled and split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()

if backend.image_data_format() == 'channels_first':
    x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
    x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
    input_shape = (1, img_rows, img_cols)
else:
    x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
    x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
    input_shape = (img_rows, img_cols, 1)

x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
x_train = random_noise(x_train, mode='speckle')
x_test = random_noise(x_test, mode='speckle')

```

```

[28]: noise_loss_2 = []

```

```

[29]: batch_size = 128
num_classes = 10

```

```

epochs = 12

# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

model = Sequential()
model.add(Conv2D(32, kernel_size=(3, 3),
                 activation='relu',
                 input_shape=input_shape))
model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))
model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))

model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Adadelta(),
              metrics=['accuracy'])

model.fit(x_train, y_train,
        batch_size=batch_size,
        epochs=epochs,
        verbose=1,
        validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
noise_loss_2.append(model.evaluate(x_test, y_test, verbose=0))
print('Test loss:', score[0])
print('Test accuracy:', score[1])

```

Train on 60000 samples, validate on 10000 samples

Epoch 1/12

60000/60000 [=====] - 59s 983us/step - loss: 0.2590 - accuracy: 0.9206 - val\_loss: 0.0593 - val\_accuracy: 0.9799

Epoch 2/12

60000/60000 [=====] - 60s 1ms/step - loss: 0.0890 - accuracy: 0.9739 - val\_loss: 0.0442 - val\_accuracy: 0.9859

Epoch 3/12

60000/60000 [=====] - 60s 1ms/step - loss: 0.0676 - accuracy: 0.9796 - val\_loss: 0.0324 - val\_accuracy: 0.9892

Epoch 4/12

60000/60000 [=====] - 60s 1ms/step - loss: 0.0556 - accuracy: 0.9837 - val\_loss: 0.0302 - val\_accuracy: 0.9899

Epoch 5/12

60000/60000 [=====] - 60s 1ms/step - loss: 0.0466 -

```

accuracy: 0.9855 - val_loss: 0.0324 - val_accuracy: 0.9890
Epoch 6/12
60000/60000 [=====] - 60s 994us/step - loss: 0.0424 -
accuracy: 0.9872 - val_loss: 0.0283 - val_accuracy: 0.9908
Epoch 7/12
60000/60000 [=====] - 60s 996us/step - loss: 0.0367 -
accuracy: 0.9886 - val_loss: 0.0262 - val_accuracy: 0.9911
Epoch 8/12
60000/60000 [=====] - 61s 1ms/step - loss: 0.0342 -
accuracy: 0.9896 - val_loss: 0.0273 - val_accuracy: 0.9913
Epoch 9/12
60000/60000 [=====] - 67s 1ms/step - loss: 0.0320 -
accuracy: 0.9904 - val_loss: 0.0304 - val_accuracy: 0.9910
Epoch 10/12
60000/60000 [=====] - 61s 1ms/step - loss: 0.0294 -
accuracy: 0.9909 - val_loss: 0.0260 - val_accuracy: 0.9918
Epoch 11/12
60000/60000 [=====] - 60s 1ms/step - loss: 0.0268 -
accuracy: 0.9916 - val_loss: 0.0243 - val_accuracy: 0.9919
Epoch 12/12
60000/60000 [=====] - 60s 1ms/step - loss: 0.0253 -
accuracy: 0.9920 - val_loss: 0.0250 - val_accuracy: 0.9926
Test loss: 0.025003947326906927
Test accuracy: 0.9926000237464905

```

```

[30]: print(convnet_loss)
      print(noise_convnet)
      print(noise_loss_2)

```

```

[[0.027638795424243655, 0.9915000200271606]]
[[0.03143596427084858, 0.9904999732971191]]
[[0.025003947326906927, 0.9926000237464905]]

```

## 2.8 Graphing the scores

```

[64]: # had to flatten the 2D array to extract the scores w.o. loop
x1 = np.array(convnet_loss).flatten()[0]
y1 = np.array(convnet_loss).flatten()[1]

x2 = np.array(noise_convnet).flatten()[0]
y2 = np.array(noise_convnet).flatten()[1]


x3 = np.array(noise_loss_2).flatten()[0]
y3 = np.array(noise_loss_2).flatten()[1]

x = [x1,x2,x3]
y = [y1,y2,y3]

```

```
plt.plot(x, y)
plt.xlabel("Loss Score for Convnet")
plt.ylabel("Accuracy Score for Convnet")
```

[64]: `Text(0,0.5,'Accuracy Score for Convnet')`



```
[70]: # flatten the lists containing loss scores.
x1 = np.array(noisey_mlenn_loss).flatten()[0]
y1 = np.array(noisey_mlenn_loss).flatten()[1]

x2 = np.array(initial_mlenn_loss).flatten()[0]
y2 = np.array(initial_mlenn_loss).flatten()[1]

x3 = np.array(more_noise_loss).flatten()[0]
y3 = np.array(more_noise_loss).flatten()[1]

x = [x2,x1,x3]
y = [y2,y1,y3]

# had to flatten the 2D array to extract the scores w.o. loop
x4 = np.array(convnet_loss).flatten()[0]
y4 = np.array(convnet_loss).flatten()[1]
```

```

x5 = np.array(noise_convnet).flatten()[0]
y5 = np.array(noise_convnet).flatten()[1]

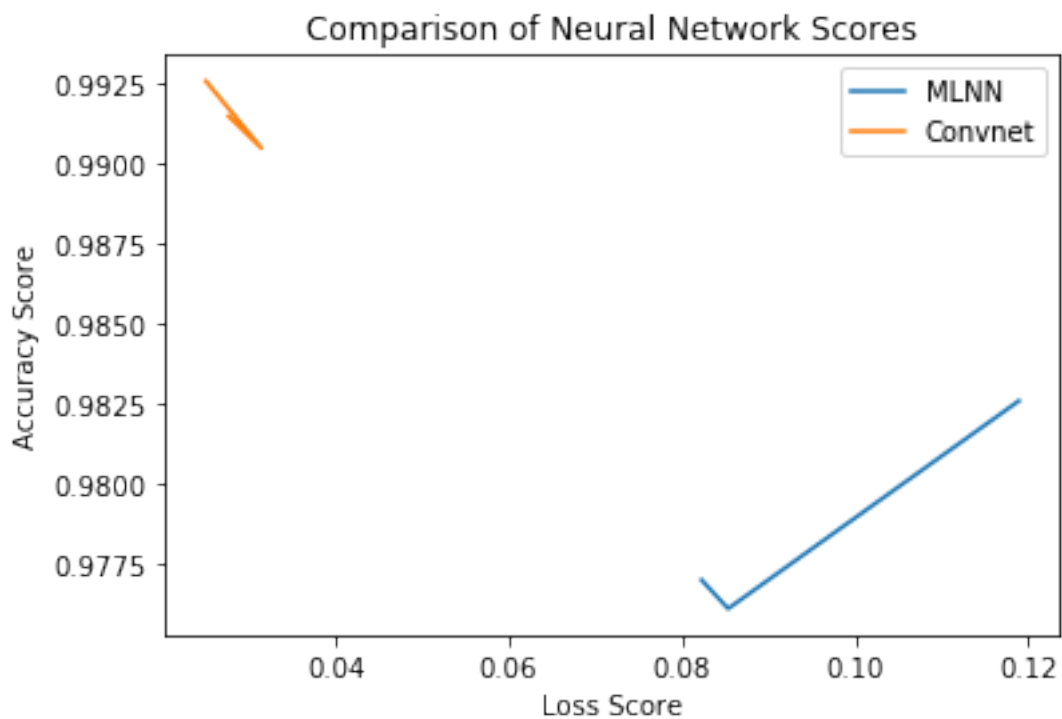
x6 = np.array(noise_loss_2).flatten()[0]
y6 = np.array(noise_loss_2).flatten()[1]

x2 = [x4,x5,x6]
y2 = [y4,y5,y6]

# graph the data.
plt.plot(x,y, label = 'MLNN')
plt.plot(x2,y2, label = 'Convnet')
plt.title("Comparison of Neural Network Scores")
plt.xlabel("Loss Score")
plt.ylabel("Accuracy Score")
plt.legend()

```

[70]: <matplotlib.legend.Legend at 0x63d757668>



[ ]: