

# Trees Regression HW 6

March 10, 2020

- <https://scikit-learn.org/stable/modules/tree.html>
- <http://scikit-learn.org/stable/modules/generated/sklearn.tree.DecisionTreeRegressor.html>

```
[1]: import sklearn
```

```
sklearn.__version__
```

```
[1]: '0.22.1'
```

```
[2]: !pip3 install scikit-learn --upgrade
```

Collecting scikit-learn

Downloading [https://files.pythonhosted.org/packages/64/57/23176044d9371e1af286176fd61cf7f74ed46d0b99122624ab93b3f32715/scikit\\_learn-0.22.2.post1-cp37-cp37m-macosx\\_10\\_9\\_x86\\_64.whl](https://files.pythonhosted.org/packages/64/57/23176044d9371e1af286176fd61cf7f74ed46d0b99122624ab93b3f32715/scikit_learn-0.22.2.post1-cp37-cp37m-macosx_10_9_x86_64.whl) (7.1MB)

100% | | 7.1MB 7.3MB/s eta 0:00:01

Collecting numpy>=1.11.0 (from scikit-learn)

Using cached [https://files.pythonhosted.org/packages/2f/5b/2cc2b9285e8b2ca8d2c1e4a2cbf1b12d70a2488ea78170de1909bca725f2/numpy-1.18.1-cp37-cp37m-macosx\\_10\\_9\\_x86\\_64.whl](https://files.pythonhosted.org/packages/2f/5b/2cc2b9285e8b2ca8d2c1e4a2cbf1b12d70a2488ea78170de1909bca725f2/numpy-1.18.1-cp37-cp37m-macosx_10_9_x86_64.whl)

Collecting scipy>=0.17.0 (from scikit-learn)

Using cached [https://files.pythonhosted.org/packages/85/7a/ae480be23b768910a9327c33517ced4623ba88dc035f9ce0206657c353a9/scipy-1.4.1-cp37-cp37m-macosx\\_10\\_6\\_intel.whl](https://files.pythonhosted.org/packages/85/7a/ae480be23b768910a9327c33517ced4623ba88dc035f9ce0206657c353a9/scipy-1.4.1-cp37-cp37m-macosx_10_6_intel.whl)

Collecting joblib>=0.11 (from scikit-learn)

Downloading <https://files.pythonhosted.org/packages/28/5c/cf6a2b65a321c4a209efcdf64c2689efae2cb62661f8f6f4bb28547cf1bf/joblib-0.14.1-py2.py3-none-any.whl> (294kB)

100% | | 296kB 28.6MB/s ta 0:00:01

Installing collected packages: numpy, scipy, joblib, scikit-learn

Could not install packages due to an EnvironmentError: [Errno 13]

Permission denied:

'/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/numpy'

Consider using the `--user` option or check the permissions.

You are using pip version 10.0.1, however version 20.0.2 is available.

You should consider upgrading via the 'pip install --upgrade pip' command.


```
[3]: import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
import pandas as pd
import numpy as np
```

```
[4]: plt.rcParams['figure.figsize'] = (20, 6)
plt.rcParams['font.size'] = 14
```

```
[5]: x = np.linspace(0, 2* np.pi, 100)
y = np.sin(x) + .5*np.random.random(100)
```

```
[6]: plt.scatter(x, y)
```

```
[6]: <matplotlib.collections.PathCollection at 0x1a184c8cc0>
```



```
[7]: from sklearn import tree
```

```
[8]: 2**16
```


```
[8]: 65536
```

```
[9]: regression = tree.DecisionTreeRegressor(max_depth=8, min_samples_split=8)
regression.fit(x.reshape(-1, 1), y)

yp = regression.predict(x.reshape(-1,1))

plt.scatter(x, yp)
plt.plot(x, y)
```

```
[9]: [<matplotlib.lines.Line2D at 0x1a1918f7b8>]
```



```
[10]: regression.predict([[2]])
```

```
[10]: array([1.34292252])
```

```
[11]: path = regression.decision_path(x.reshape(-1, 1))
```

```
[12]: path.todense()
```

```
[12]: matrix([[1, 1, 1, ..., 0, 0, 0],
             [1, 1, 1, ..., 0, 0, 0],
             [1, 1, 1, ..., 0, 0, 0],
             ...,
             [1, 0, 0, ..., 1, 0, 1],
             [1, 0, 0, ..., 1, 0, 1],
             [1, 0, 0, ..., 1, 0, 1]])
```

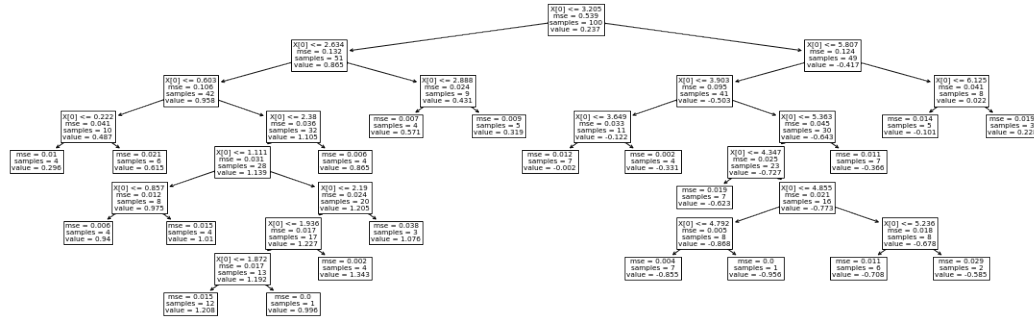
```
[13]: tree.plot_tree(regression)
```

```
[13]: [Text(614.976,314.651,'X[0] <= 3.205\nmse = 0.539\nsamples = 100\nvalue =
0.237'),
      Text(347.595,277.633,'X[0] <= 2.634\nmse = 0.132\nsamples = 51\nvalue =
0.865'),
      Text(213.905,240.616,'X[0] <= 0.603\nmse = 0.106\nsamples = 42\nvalue =
0.958'),
      Text(106.952,203.598,'X[0] <= 0.222\nmse = 0.041\nsamples = 10\nvalue =
0.487'),
      Text(53.4762,166.58,'mse = 0.01\nsamples = 4\nvalue = 0.296'),
      Text(160.429,166.58,'mse = 0.021\nsamples = 6\nvalue = 0.615'),
      Text(320.857,203.598,'X[0] <= 2.38\nmse = 0.036\nsamples = 32\nvalue = 1.105'),
      Text(267.381,166.58,'X[0] <= 1.111\nmse = 0.031\nsamples = 28\nvalue = 1.139'),
      Text(160.429,129.562,'X[0] <= 0.857\nmse = 0.012\nsamples = 8\nvalue = 0.975'),
      Text(106.952,92.5444,'mse = 0.006\nsamples = 4\nvalue = 0.94'),
      Text(213.905,92.5444,'mse = 0.015\nsamples = 4\nvalue = 1.01'),
      Text(374.333,129.562,'X[0] <= 2.19\nmse = 0.024\nsamples = 20\nvalue = 1.205'),
```

```

Text(320.857,92.5444,'X[0] <= 1.936\nmse = 0.017\nsamples = 17\nvalue =
1.227'),
Text(267.381,55.5267,'X[0] <= 1.872\nmse = 0.017\nsamples = 13\nvalue =
1.192'),
Text(213.905,18.5089,'mse = 0.015\nsamples = 12\nvalue = 1.208'),
Text(320.857,18.5089,'mse = 0.0\nsamples = 1\nvalue = 0.996'),
Text(374.333,55.5267,'mse = 0.002\nsamples = 4\nvalue = 1.343'),
Text(427.81,92.5444,'mse = 0.038\nsamples = 3\nvalue = 1.076'),
Text(374.333,166.58,'mse = 0.006\nsamples = 4\nvalue = 0.865'),
Text(481.286,240.616,'X[0] <= 2.888\nmse = 0.024\nsamples = 9\nvalue = 0.431'),
Text(427.81,203.598,'mse = 0.007\nsamples = 4\nvalue = 0.571'),
Text(534.762,203.598,'mse = 0.009\nsamples = 5\nvalue = 0.319'),
Text(882.357,277.633,'X[0] <= 5.807\nmse = 0.124\nsamples = 49\nvalue =
-0.417'),
Text(748.667,240.616,'X[0] <= 3.903\nmse = 0.095\nsamples = 41\nvalue =
-0.503'),
Text(641.714,203.598,'X[0] <= 3.649\nmse = 0.033\nsamples = 11\nvalue =
-0.122'),
Text(588.238,166.58,'mse = 0.012\nsamples = 7\nvalue = -0.002'),
Text(695.19,166.58,'mse = 0.002\nsamples = 4\nvalue = -0.331'),
Text(855.619,203.598,'X[0] <= 5.363\nmse = 0.045\nsamples = 30\nvalue =
-0.643'),
Text(802.143,166.58,'X[0] <= 4.347\nmse = 0.025\nsamples = 23\nvalue =
-0.727'),
Text(748.667,129.562,'mse = 0.019\nsamples = 7\nvalue = -0.623'),
Text(855.619,129.562,'X[0] <= 4.855\nmse = 0.021\nsamples = 16\nvalue =
-0.773'),
Text(748.667,92.5444,'X[0] <= 4.792\nmse = 0.005\nsamples = 8\nvalue =
-0.868'),
Text(695.19,55.5267,'mse = 0.004\nsamples = 7\nvalue = -0.855'),
Text(802.143,55.5267,'mse = 0.0\nsamples = 1\nvalue = -0.956'),
Text(962.571,92.5444,'X[0] <= 5.236\nmse = 0.018\nsamples = 8\nvalue =
-0.678'),
Text(909.095,55.5267,'mse = 0.011\nsamples = 6\nvalue = -0.708'),
Text(1016.05,55.5267,'mse = 0.029\nsamples = 2\nvalue = -0.585'),
Text(909.095,166.58,'mse = 0.011\nsamples = 7\nvalue = -0.366'),
Text(1016.05,240.616,'X[0] <= 6.125\nmse = 0.041\nsamples = 8\nvalue = 0.022'),
Text(962.571,203.598,'mse = 0.014\nsamples = 5\nvalue = -0.101'),
Text(1069.52,203.598,'mse = 0.019\nsamples = 3\nvalue = 0.228')]

```



```
[14]: bikeshare = pd.read_csv('../data/bikeshare_daily_agg.csv',
    ↪index_col='hour_of_day')
```

```
[15]: bikeshare
```


```
[15]:
```

	0	1	2	3	4	5	6
hour_of_day							
0.0	21.0	34.0	43.0	47.0	51.0	89.0	106.0
0.1	39.0	22.0	27.0	37.0	56.0	87.0	100.0
0.2	31.0	24.0	26.0	42.0	50.0	98.0	77.0
0.3	26.0	27.0	25.0	29.0	52.0	99.0	87.0
0.4	19.0	24.0	29.0	29.0	50.0	98.0	69.0
...	...	...	...	...	...	...	...
23.5	36.0	65.0	60.0	94.0	80.0	93.0	28.0
23.6	37.0	61.0	66.0	100.0	81.0	95.0	28.0
23.7	30.0	42.0	49.0	80.0	101.0	105.0	27.0
23.8	33.0	52.0	47.0	79.0	91.0	93.0	24.0
23.9	34.0	33.0	48.0	65.0	105.0	111.0	23.0

[240 rows x 7 columns]


```
[16]: plt.plot(bikeshare)
```

```
[16]: [<matplotlib.lines.Line2D at 0x103b2b550>,
    <matplotlib.lines.Line2D at 0x103b2b748>,
    <matplotlib.lines.Line2D at 0x103b2b898>,
    <matplotlib.lines.Line2D at 0x103b2b9e8>,
    <matplotlib.lines.Line2D at 0x103b2bb38>,
    <matplotlib.lines.Line2D at 0x103b2bc88>,
    <matplotlib.lines.Line2D at 0x103b2bdd8>]
```



```
[17]: plt.plot(bikeshare['5'])
```

```
[17]: [<matplotlib.lines.Line2D at 0x1a19a45860>]
```




```
[18]: hours = bikeshare.index.values.reshape(-1,1)

bike_reg = tree.DecisionTreeRegressor(max_depth=5)
bike_reg.fit(hours, bikeshare['5'].fillna(0))

bike_pred = bike_reg.predict(hours)

plt.scatter(hours, bike_pred, s=2)
plt.scatter(hours, bikeshare['5'], s=2)
```

```
[18]: <matplotlib.collections.PathCollection at 0x1a19c76198>
```



- 1 Use the bikeshare dataset (see above) and choose a weekday (0,1,2,3,4).
- 2 1. Create 5 Decision Tree Regressors using `max_depth=4,5,6,7,8`. For each one of these models, calculate the MSE between the predicted values from the model (`bike_pred`) and the actual values (`bikeshare['n']`). Create a plot showing the predictions along with the actuals. You may also show the `print_tree()` for a sanity check as well.
- 3 2. Using the 5 models created with various `max_depth` values, calculate the MSE between the predicted values (`bike_pred`) and values from all of the weekdays [0,1,2,3,4]. You should have 25 total MSE values, 5 values for each `max_depth`.
- 4 3. (2 cont'd) Describe which `max_depth` you would recommend based on the groups of MSE values. Use the idea of generality of the model for your argument along with the MSE values as proof.

```
[167]: from sklearn.metrics import mean_squared_error
from sklearn.model_selection import train_test_split

# extract subset of bikeshare data for monday and organize the dataframe.
monday_data = bikeshare.loc[:, '0'].reset_index().fillna(0)
print(monday_data.shape)
monday_data.columns = ['hour', 'distance']

# peace of mind that theres no NaN values in the data set.
```

```

print('Total NaN vals: ' + str(monday_data.distance.isna().sum())) # verify the
    ↳correction worked.

# test and train data.
x_train, x_test, y_train, y_test = train_test_split(monday_data.hour,
    ↳monday_data.distance, train_size = .8)

# iterate through 4:8 for max depth.
for i in range(4,9):

    # create model to answer Q1.
    regression = tree.DecisionTreeRegressor(max_depth=i)
    regression.fit(np.array(x_train).reshape(-1,1) , y_train)
    # predict values on model.
    pred_vals = regression.predict(np.array(monday_data.hour).reshape(-1,1))

    # plot graphs onto the original data.
    plt.scatter(monday_data.hour, pred_vals)
    plt.legend()
    plt.plot(monday_data.hour, monday_data.distance)
    mse = mean_squared_error(monday_data.distance, pred_vals)
    print('For max_depth = ' + str(i) + ": MSE = " + str(mse))

    for j in range(0,5):
        day_of_week = bikeshare.iloc[:,j].reset_index().fillna(0)
        #print(day_of_week)
        #print(day_of_week.iloc[:,1])
        mse_per_day = mean_squared_error(bikeshare.iloc[:,j].fillna(0),
    ↳pred_vals)

        print(f'For day = {j} MSE = {mse_per_day:0.2f} and max_depth = {i}
    ↳RMSE = {np.sqrt(mse_per_day):0.2f}')
    # run model against the other days of the week in the data set.

```

No handles with labels found to put in legend.

(240, 2)

Total NaN vals: 0

For max\_depth = 4: MSE = 10376.405722311105

For day = 0 MSE = 10376.41 and max\_depth = 4 RMSE = 101.86

For day = 1 MSE = 14325.53 and max\_depth = 4 RMSE = 119.69

For day = 2 MSE = 19741.58 and max\_depth = 4 RMSE = 140.50


For day = 3 MSE = 28348.12 and max\_depth = 4 RMSE = 168.37



```

For day = 4 MSE = 18790.31 and max_depth = 4 RMSE = 137.08
For max_depth = 5: MSE = 3904.3691496921188
For day = 0 MSE = 3904.37 and max_depth = 5 RMSE = 62.48
For day = 1 MSE = 7368.86 and max_depth = 5 RMSE = 85.84
For day = 2 MSE = 12060.61 and max_depth = 5 RMSE = 109.82
For day = 3 MSE = 19023.68 and max_depth = 5 RMSE = 137.93
For day = 4 MSE = 11164.04 and max_depth = 5 RMSE = 105.66
For max_depth = 6: MSE = 1682.6231752245399
For day = 0 MSE = 1682.62 and max_depth = 6 RMSE = 41.02
For day = 1 MSE = 4276.64 and max_depth = 6 RMSE = 65.40
For day = 2 MSE = 8153.59 and max_depth = 6 RMSE = 90.30
For day = 3 MSE = 14729.12 and max_depth = 6 RMSE = 121.36
For day = 4 MSE = 8890.57 and max_depth = 6 RMSE = 94.29
For max_depth = 7: MSE = 661.8262039410333
For day = 0 MSE = 661.83 and max_depth = 7 RMSE = 25.73
For day = 1 MSE = 3081.19 and max_depth = 7 RMSE = 55.51
For day = 2 MSE = 6776.42 and max_depth = 7 RMSE = 82.32
For day = 3 MSE = 12661.36 and max_depth = 7 RMSE = 112.52
For day = 4 MSE = 6973.01 and max_depth = 7 RMSE = 83.50
For max_depth = 8: MSE = 406.88504399061736
For day = 0 MSE = 406.89 and max_depth = 8 RMSE = 20.17
For day = 1 MSE = 2568.42 and max_depth = 8 RMSE = 50.68
For day = 2 MSE = 6064.12 and max_depth = 8 RMSE = 77.87
For day = 3 MSE = 11996.04 and max_depth = 8 RMSE = 109.53
For day = 4 MSE = 6356.28 and max_depth = 8 RMSE = 79.73

```



For question 3, I would suggest using a max\_depth score of 8 since the avg RMSE is 67.59 while for max\_depth = 7 mean RMSE is 71.91, max\_depth = 6 mean RMSE is 82.47, max\_depth = 5 mean RMSE is 100.35, max\_depth = 4 mean RMSE is 133.5

[ ]: